

Candidate Ranking - Business Recommendations

Final approach: NLPL Model 40 Word2Vec + transparent relevance rules + recruiter feedback

0.948 NDCG@10 - aspiring HR	0.935 NDCG@10 - seeking HR	100% meaningful-token coverage	104 -> 52 rows -> unique titles
---------------------------------------	--------------------------------------	--	--

What was tested

The two recruiter searches were **aspiring human resources** and **seeking human resources**. Duplicate job titles were collapsed before evaluation, leaving 52 unique titles. A transparent relevance score, $R = H(0.70 + 0.30I)$, was used as the graded benchmark; Model 40 then ranked titles by semantic similarity.

Word2Vec

Word2Vec represents words as numerical vectors learned from language usage, allowing titles with related meaning to be compared even when wording differs.

NDCG@10

NDCG@10 measures how well the ten highest-ranked results put the most relevant candidates nearest the top; 1.0 is ideal.

What the results mean

Model 40 produced strong rankings for both searches and covered every meaningful token used in the titles and queries. This gives the ranking a solid semantic base without relying on location or connection count as shortcuts for candidate quality.

Recruiter feedback is a controlled personalization dial. At 30% influence, both representative starred candidates moved from #2 to #1 while the original top-ten retained 9/10 and 10/10 candidates. Across all 34 strong scenarios, stronger feedback moves the starred candidate upward faster but changes more of the original shortlist, so there is no single universal best weight.

Recommendation. Use the rule score to flag clearly irrelevant profiles, Model 40 to order relevant candidates, and recruiter feedback to personalize the shortlist. Keep 30% as a conservative starting point, but do not market it as universally optimal. Log recruiter actions and tune feedback strength from real usage.